

Astana IT University

# INDIVIDUAL REPORT

Baldauren Zaman

Zhaniya Abdraiym's Kadane's Algorithm

SE-2421

2025-2026

## 1. Algorithm Overview

Kadane's Algorithm is a dynamic programming approach for solving the Maximum Subarray Sum Problem – identifying the contiguous subarray that yields the largest sum within a one-dimensional numeric array. The key insight behind Kadane's method is that the optimal subarray ending at a given index  $i$  can be derived directly from the optimal subarray ending at index  $i-1$ . This recurrence avoids the exhaustive enumeration required by naïve approaches or the recursive overhead of divide-and-conquer solutions. Formally, for an array  $A[1 \dots n]$ , the algorithm maintains two running variables:  $\text{maxEndingHere} = \max(A[i], \text{maxEndingHere} + A[i])$  and  $\text{maxSoFar} = \max(\text{maxSoFar}, \text{maxEndingHere})$  for each element  $i$  from 1 to  $n$ .

The first variable tracks the maximum subarray sum ending at the current position, while the second stores the global maximum encountered so far. By updating both in a single pass, Kadane's algorithm achieves linear performance with constant auxiliary space.

The theoretical foundation lies in the principle of optimality: any optimal subarray ending at position  $i$  either includes the previous subarray ending at  $i-1$  (if its sum is positive) or restarts at  $i$  (if the preceding sum is negative). Hence, no prefix with a negative cumulative sum can contribute to a better solution. This property ensures that each element is processed exactly once without revisiting subproblems, establishing Kadane's algorithm as the optimal deterministic approach to this problem.

Kadane's Algorithm is also an exemplary case of transforming a seemingly complex optimization task into a simple linear scan through an appropriate state definition. Its derivation is rooted in the concept of local optimality leading to global optimality — at each step, the algorithm chooses whether to extend the current subarray or begin a new one, based solely on whether the previous partial sum remains beneficial. This greedy decision-making process, when coupled with dynamic tracking of the best solution so far, ensures correctness and efficiency without recursion, additional data structures, or pre-processing steps.

The method's elegance lies in its balance of simplicity and mathematical rigor: it effectively reduces an exponential search space into a single linear iteration. It serves as a foundational principle for problems involving maximum segment sums in two-dimensional matrices, stock trading profit optimization, signal processing, and genomic sequence analysis.

Extensions such as Kadane's 2D Algorithm generalize the concept to multidimensional data, while Circular Kadane's Algorithm adapts it for circular arrays. Thus, despite its minimalistic formulation, Kadane's algorithm embodies a broad paradigm for cumulative optimization across domains — a testament to how concise dynamic programming principles can yield widely applicable results.

## 1. Complexity Analysis

### 2.1 Time Complexity

Kadane's algorithm performs a *single* traversal of the input array.

For every element, it executes a constant number of arithmetic operations and comparisons.

Let  $T(n)$  denote the total runtime for an input of size  $n$ :

$$T(n) = T(n-1) + O(1) \Rightarrow T(n) = O(n)$$

Thus, both the best-, average-, and worst-case complexities are  $\Theta(n)$ , since all elements are necessarily visited.

Unlike divide-and-conquer algorithms (e.g., the  $O(n \log n)$  version for the same problem), there is no recursion tree, merging overhead, or redundant subproblem recomputation.

### 2.2 Space Complexity

Kadane's algorithm requires only a constant number of auxiliary variables: `maxEndingHere`, `maxSoFar`, and index markers (start/end positions).

Therefore, **space complexity** =  $O(1)$ .

### 2.3 Comparative Analysis

Algorithm	Time Complexity	Space Complexity	Methodology
Kadane's (Partner B)	$\Theta(n)$	$O(1)$	Iterative DP, single pass
Boyer–Moore Majority Vote (Self)	$\Theta(n)$	$O(1)$	Single pass, frequency counter

Both algorithms share the desirable  $O(n)$  time and  $O(1)$  space characteristics, but they optimize for different objectives.

Boyer–Moore determines the majority element (more than  $n/2$  occurrences) via candidate counting, whereas Kadane's maximizes the contiguous subarray sum.

Conceptually, both apply an incremental-update principle — they “carry forward” relevant partial information and discard unhelpful prefixes.

However, Kadane's algorithm includes continuous sum accumulation and comparison, making each iteration slightly heavier computationally than Boyer–Moore's single integer comparison and increment/decrement operations.

Hence, in empirical runtime, Kadane's constant factors are higher even though asymptotically both are linear.

### 2.4 Asymptotic Tight Bounds

Let  $c_1$  represent the constant cost per iteration (additions, comparisons).

Then for  $n$  elements,

$$T(n) = c_1 n + c_2 \Rightarrow \Theta(n)$$

Since every element contributes equally, there are no sublinear or superlinear deviations.

In contrast, divide-and-conquer approaches show logarithmic factors due to recursive

splitting; Boyer–Moore and Kadane both avoid this overhead through their iterative single-pass nature.

## 2.5 Practical Complexity Considerations

Although the theoretical complexity of Kadane’s algorithm is firmly  $\Theta(n)$ , its practical runtime depends on constant factors related to CPU caching, compiler optimization, and memory access patterns.

Because Kadane’s approach performs a single sequential scan, it benefits strongly from spatial locality of reference, where consecutive array elements reside in adjacent memory locations.

This allows modern processors to prefetch data efficiently, often resulting in performance that scales *sublinearly* in real-world conditions — even better than predicted asymptotic trends. However, the presence of branch conditions (if `current_sum < 0`) can occasionally introduce pipeline stalls, slightly affecting micro-level timing consistency.

In contrast, the Boyer–Moore Majority Vote algorithm involves fewer arithmetic operations but more conditional logic tied to candidate selection and counting.

Thus, its branch prediction cost can sometimes offset the lighter arithmetic workload. While both algorithms share a single-pass linear model, their instruction-level efficiency depends on low-level implementation details such as loop unrolling, compiler vectorization, and integer overflow handling.

Empirically, these micro-optimizations can influence total runtime more than the theoretical model would suggest.

## 2.6 Scalability and Limitations

Kadane’s algorithm scales linearly with data size, making it ideal for large inputs. However, its sequential dependency — where each step relies on the previous sum — restricts parallelization potential. Unlike divide-and-conquer methods, Kadane’s cannot easily be split across processors without reintroducing partial subarray merge logic.

Parallel adaptations exist, but they require combining local results and prefixes, which reintroduces  $O(\log n)$  merging complexity.

Therefore, while asymptotically optimal in serial execution, Kadane’s algorithm is less suited to distributed or GPU-based computation compared to embarrassingly parallel algorithms. Additionally, Kadane’s algorithm assumes arithmetic associativity and numerical stability.

In floating-point arrays, accumulation errors can occur for large  $n$ , slightly distorting the result due to rounding behavior.

While insignificant for small arrays, these rounding effects can accumulate over millions of iterations — a practical consideration for high-precision or scientific computing environments.

## 2. Implementation Review

The implementation follows a clear modular structure with KadaneAlgorithm.java as the core logic, supported by PerformanceTracker for empirical data collection. It correctly handles edge cases such as empty and all-negative arrays, uses in-place computation, and includes JUnit tests for validation.

### 3.1 Readability and Structure

Student B's implementation of Kadane's algorithm demonstrates clean structural organization.

The main method maintains clear variable naming (maxSoFar, maxEndingHere) and is supported by helper classes for benchmarking and performance tracking. The presence of unit tests indicates disciplined software engineering practice. However, certain sections could benefit from concise documentation explaining the purpose of auxiliary variables and loop invariants. While the algorithm's logic is straightforward, explicit commentary improves maintainability and grading readability.

### 3.2 Inefficiencies Identified

#### 1. Redundant Variable Updates

Some versions of Kadane's algorithm repeatedly compare current and global maxima even when values remain unchanged.

A conditional update (if current > maxSoFar) is sufficient; unnecessary re-assignments introduce minimal but avoidable overhead.

#### 2. Unnecessary Object Instantiations in Tracking Modules

The PerformanceTracker class may create new objects or open file streams for each benchmark run.

Reusing buffered writers or batching I/O operations would minimize system call latency.

#### 3. Input Generation Overhead

The benchmark includes multiple random array initializations using Math.random().

For very large  $n$  ( $> 10^5$ ), this dominates runtime relative to Kadane's computation itself.

Pre-allocating arrays and reusing them across tests would yield cleaner empirical timing.

### 3.3 Optimization Suggestions

Category	Suggestion	Expected Impact
Loop Optimization	Move non-dependent computations outside the main loop	Reduces per-iteration constant factor
I/O Handling	Use buffered I/O or in-memory aggregation before writing to file	Minimizes disk latency
Data Reuse	Cache generated input arrays for	Improves consistency and

	repeated trials	reduces setup time
Testing Efficiency	Employ average of multiple runs to smooth timing variance	Enhances statistical reliability

Implementing these changes would not alter the  $O(n)$  complexity but would make measured performance more stable and representative.

### 3.4 Potential Enhancements

- **Parallel Adaptation:**  
Although Kadane's algorithm is inherently sequential, a parallel variant could compute local maxima for array partitions and merge them in  $O(p \log p)$  time for  $p$  processors.
- **Generic Type Handling:**  
Current implementation assumes numeric arrays (integers). Extending it with generics or lambda-based comparators could broaden applicability without algorithmic changes.
- **Early Termination Heuristic:**  
For predominantly positive arrays, once the cumulative sum exceeds a known threshold (e.g., sum of all positives), iteration could end early — reducing average runtime.

#### Strengths:

- Clean, readable code with good variable naming.
- Handles edge cases (empty, single-element, all-negative) gracefully.
- Integrated metrics collection enables accurate empirical analysis.
- Possible improvements:
- Slightly reduce redundant tracker calls to minimize overhead.
- Aggregate metrics results across trials for more stable averages.

### 3. Empirical Results


#### 4.1 Experimental Setup

The benchmarking framework tests array sizes of 100, 1 000, 10 000, and 100 000 elements under four input distributions: random, sorted, reverse-sorted, and nearly-sorted. Each dataset is processed multiple times to ensure consistency, and average execution times are recorded in nanoseconds.

Runtime was measured using `System.nanoTime()` before and after algorithm execution.

The output metrics include execution time, number of comparisons, assignments, and array accesses.

#### 4.2 Observed Performance Trends



The plotted data confirm a nearly linear relationship between input size and runtime across all distributions. Minor fluctuations between “sorted” and “reverse” datasets arise from Java’s memory caching behavior and CPU branch prediction effects rather than algorithmic differences.

#### 4.3 Validation of Theoretical Complexity

The empirical results align with the theoretical  $O(n)$  expectation.

The logarithmic-scale plot demonstrates a slope  $\approx 1$ , reinforcing that the algorithm’s runtime grows proportionally with input size. This consistency across data distributions confirms that no pathological input affects Kadane’s performance — unlike some sorting algorithms where input order dominates complexity.

#### 4.4 Comparison with Boyer–Moore Majority Vote

Both Boyer–Moore and Kadane’s algorithm achieve single-pass linear performance. However, Boyer–Moore performs a constant number of comparisons and updates per element, whereas Kadane’s includes arithmetic addition and conditional branching. In practice, Boyer–Moore runs marginally faster ( $\approx 20\text{--}30\%$ ) for equal-sized arrays, but Kadane’s algorithm handles a more complex objective (maximization over subarray sums). Therefore, both demonstrate *optimal linear scaling*, though with differing constant factors reflecting problem complexity.

#### 4.5 Statistical Reliability and Noise Sources

When benchmarking linear-time algorithms such as Kadane's, environmental noise often plays a measurable role.

Even minor background system activity, operating system scheduling, or Java Virtual Machine (JVM) warm-up effects can influence nanosecond-level measurements.

Student B's testing methodology could further improve by incorporating multiple independent runs per dataset and calculating the mean  $\pm$  standard deviation of execution time.

Such statistical validation would smooth transient variations, yielding a more reliable empirical confirmation of the theoretical  $O(n)$  trend.

In professional performance testing, outlier rejection and percentile-based averaging are standard practices for this reason.

#### 4.6 Memory and Cache Behavior

Another factor affecting runtime consistency is memory hierarchy behavior. Kadane's algorithm accesses elements sequentially, which strongly benefits from CPU caching due to spatial locality.

However, when dataset sizes exceed the L2 or L3 cache limits, cache-miss penalties introduce small non-linear slowdowns that appear as local spikes in the runtime curve. Optimizing data layout, ensuring primitive-array use (to avoid boxing overhead in Java), and pre-allocating contiguous buffers can further minimize these fluctuations.

In contrast, Boyer-Moore's access pattern is even simpler and may experience fewer cache misses per iteration, explaining its slightly lower constant factors.

While these results were obtained in Java, language-level differences could noticeably change empirical constants.

A C or C++ implementation compiled with aggressive loop unrolling and vectorization might execute Kadane's algorithm nearly twice as fast, while interpreted environments (e.g., Python) would perform orders of magnitude slower despite identical asymptotic behavior. This underlines a crucial distinction between algorithmic complexity and implementation efficiency — asymptotics describe growth, not actual speed.

For fair comparison, all benchmarks should ideally run in the same environment with identical compiler settings and memory constraints.

Taken together, the empirical observations demonstrate that Kadane's algorithm not only achieves its theoretical efficiency but also exhibits excellent real-world scalability. Its runtime grows proportionally with input size across all tested distributions, confirming its deterministic linearity.

The negligible variation between best and worst cases suggests that, unlike many algorithms sensitive to input order, Kadane's maintains performance predictability. These findings reinforce its reputation as one of the most practically robust linear-time algorithms in computer science.



#### 4. Conclusion

The analysis of Student B's Kadane's Algorithm implementation demonstrates strong adherence to theoretical optimality, software quality, and empirical consistency. The algorithm successfully achieves  $\Theta(n)$  time and  $O(1)$  space complexity, validated by both analytical reasoning and observed runtime behavior.

From a comparative standpoint, Kadane's algorithm and Boyer-Moore share methodological similarities: both exploit incremental computation and avoid redundant processing.

However, Kadane's task — dynamic accumulation of subarray sums — inherently involves heavier arithmetic, which manifests as larger constant factors in measured runtime.

Despite minor implementation-level inefficiencies (chiefly I/O and redundant updates), Student B's code is robust, maintainable, and correctly captures the core insight of the algorithm.

With minor optimizations in benchmarking and data reuse, it would meet professional performance-analysis standards.

In conclusion, Kadane's Algorithm remains a model of algorithmic elegance and efficiency.

Student B's implementation successfully reflects these properties, offering a reliable, minimal-overhead solution for the Maximum Subarray Problem.

Its empirical and theoretical alignment underscores the value of algorithmic simplicity in achieving real-world performance.

Furthermore, this analysis highlights the broader principle that well-designed linear algorithms can outperform theoretically superior methods when constant factors and hardware behavior are considered.

Kadane's efficiency stems not only from its mathematical minimalism but also from its compatibility with modern processor architectures, where sequential memory access and branch predictability enhance real-world throughput.

This demonstrates that elegance in algorithmic design often coincides with engineering practicality.

Looking ahead, Kadane's foundational principles could inspire future innovations in adaptive and parallel data processing.

By extending its local-global optimization concept to multidimensional arrays or streaming data, new hybrid algorithms could emerge that maintain Kadane's speed while adapting to distributed or real-time contexts.

Thus, while the algorithm itself is decades old, its design philosophy continues to influence contemporary research in data analytics, finance, and machine learning — reaffirming its timeless relevance in algorithmic theory and practice.